

Agentic AI Applications

From Concept to Code

Agentic AI

Module 2

Today's activities



- Agents reasoning and planning
- Building and customizing agents
- Building AI agents responsibly
- AI agents with guardrails
- Evaluating AI Agents

How to use this slide deck depending on your persona

This full slide deck can be split according to the following two personas:

 code-based faculty

 code-free faculty

- At the bottom left of each slide you can find the icons that apply for the respective personas
- You can use them as a guide to select the proper slides according to your needs
- The slides pointing to hands-on labs also have the same icons to guide you to select the ones that best fit to your faculty's persona

Possible combinations

- Slides for code-free personas: 
- Slides for code-based personas: 
- Slide for all personas:  



Code-free personas can skip the code-based slides without compromising the full understanding of the concepts covered


Reasoning and Planning

Reasoning in LLMs


Success of traditional LLMs

- Demonstrated remarkable capabilities in several tasks:
 - Zero-shot application
 - In-context techniques


A few examples




Text
Generation



Chatbots



Classification



Translation



Summarization



Are LLMs just as good at everything?

- LLMs perform well based on:
 - Pattern recognition
 - Conceptual correlation
- Primary training objective is driven by predicting the next word
 - NOT understanding the underlying concept or logic
- Often fails when requiring genuine logical deduction
 - Core aspect of human intelligence



LLMs can fail on simple arithmetic tasks...

Counting the number of letters

Counting

Count the letters.

Input 1: iiiiiiiiiiiiiiiiiiiiiiiiiiiiii

Correct: 30

✓ **GPT-4:** 30

Input 2: iiiiiiiiiiiiiiiiiiiiiiiiiiiiii

Correct: 29

✗ **GPT-4:** 30

Simple arithmetic problems

Linear functions

Multiply by 9/5 and add 32.

Input: 328

Correct: 622.4

✓ **GPT-4:** 622.4

Multiply by 7/5 and add 31.

Input: 328

Correct: 490.2

✗ **GPT-4:** 457.6

Source: arxiv.org/abs/2309.13638



LLMs can fail on text manipulation tasks...

Swapping words

Article swapping

Swap each article (*a*, *an*, or *the*) with the word before it.

Input 1: It does not specify time a limit for registration the procedures.

Correct: It does not specify a time limit for the registration procedures.

✓ **GPT-4:** It does not specify a time limit for the registration procedures.

Input 2: It few with it to lying take the get just a hands would kinds.

Correct: It few with it to lying the take get a just hands would kinds.

✗ **GPT-4:** It flew with a few kinds to take the lying just to get the hands.

Shift ciphers

Decode by shifting each letter 13 positions backward in the alphabet.

Input: Jryy, vg jnf abg rknpgyl cynaarq sebz gur ortvaavat.

Correct: Well, it was not exactly planned from the beginning.

✓ **GPT-4:** Well, it was not exactly planned from the beginning.

Decode by shifting each letter 12 positions backward in the alphabet.

Input: Iqxx, uf ime zaf qjmofxk bxmzzqp rday ftq nqsuzzuzs.

Correct: Well, it was not exactly planned from the beginning.

✗ **GPT-4:** Wait, we are not prepared for the apocalypse yet.

Shifting each letter

Source: arxiv.org/abs/2309.13638



System 1 and System 2 thinking

- ⚙ Concepts from **cognitive science**

- » Inspiration for a lot of research in reasoning and agentic AI

*The automatic operations of **System 1** generate surprisingly complex patterns of ideas, but only the slower **System 2** can construct thoughts in an orderly series of steps.*


– **Daniel Kahneman**
Thinking, Fast and Slow



Characteristics of System 1 and System 2 thinking



Why are you
so slow?



Why are you
so stupid?


System 1	System 2
Fast	Slow
Reactive	Deliberative
Error prone	Reliable
Suitable for trivial, routine tasks	Suitable for complex decisions and planning



Which approach should you use?



- System 1:
 - Common knowledge responses
 - Such as content generation and conversational tasks
 - Direct pattern matching
 - Quick transformations
 - Trivial heuristic decisions



- System 2:
 - Strong logical reasoning (such as science, arithmetic and coding)
 - Thorough planning and execution
 - Complex orchestration workflows



Agents need to reason



Towards improving reasoning

- 
- **Implicit reasoning from traditional pre-training and post-training**

- 
- **Surface reasoning using prompting techniques**

- 
- **Reasoning models**




Reasoning and Planning

Reasoning and Planning Techniques

Chain-of-thought (CoT) prompting

- Break down complex tasks and then answer
 - Smaller, logical goals
 - Easy to solve goals




Source: [Wei et al. 2023](#)




Zero-shot chain-of-thought

- Zero-shot CoT
 - Ask the model to think step-by-step
- You can also engage CoT using examples
 - Few-shot CoT



Standard Prompting



Chain-of-thought Prompting

Source: [Wei et al. 2023](#)




Plan-and-execute

- Plan first, execute later strategy

Source: [Wang et al. 2023](#)

- Two phases:
 - **Planning phase**: Analyze task objectives, decompose into subtasks, build execution plan
 - **Execution phase**: Execute planned subtasks in sequence, process outputs



ReAct

- Interleave reasoning traces and task-specific actions
- Core technique used by most **agentic frameworks**
 - Interact with external tools to retrieve information or take actions
- Allows an LLM access to the external environment
 - Typically implemented as a cyclic process
 - Think one step at a time
 - **Dynamically** create and adjust plans


Source: [Yao et al. 2023](#)



ReAct

ReAct: Reason + Act

- LLMs are capable of reasoning and generating **action** plans.
- Combine both!
 - **Reasoning**: Create, track, and update action plans. Handle errors.
 - **Acting**: Interface with functions, tools, knowledge bases, or environments




ReAct vs Plan-and-execute

ReAct	Plan-and-execute
Dynamically alternate between thinking and acting	Plan first, execute later
Update plans more frequently based on execution	Update plans only after entire plan is executed
Typically, faster to generating final answer	Typically, slower in generating final answer
Fewer tokens consumed/generated	More tokens consumed generated
Suited for medium difficulty tasks	Suited for higher difficulty tasks



Reflection

- Prompting techniques for LLMs to reflect on and critique past actions



Other specialized reflection techniques

- Reflexion ([Shinn et al. 2023](#))
 - Critique is grounded in external data or tool use
 - Generate citations and highlight superfluous or missing content
 - More constructive reflection!
- Language agent tree search ([Zhou et al. 2024](#))
 - Combine reflection, evaluation and search for better performance
 - RL-based framing for adaptation and complex problem solving
 - Avoid getting stuck in repetitive loops




Reasoning and Planning

Reasoning Models

Era of reasoning models

- Wave of new self-reflective models



What are reasoning models?

- Traditional LLMs are forced to answer immediately
- Complex problems need more ‘**thought**’
 - Arithmetic problems, reasoning problems, etc.
- **Train models** to think critically before generating responses
- Models generate **reasoning/thought tokens** before responding.




Chat models vs reasoning models

	Traditional chat models	Reasoning models
Inference	Fast and intuitive	Slow but effortful
Reasoning	Basic chain-of-thought when prompted	High quality chain-of-thought by default
Performance	Lower benchmark performance on tasks such as math and coding problems	Higher benchmark performance on reasoning tasks
Cost	Typically, cheaper prompting costs	Typically higher prompting costs



Reasoning models in multi-turn conversations



Why do reasoning models perform so well?

- Natively work through problems inspired by how a human would
- Reasoning tokens allow creating rich context and analysis
- Number of turns are variable
 - Decided by the model or set manually
- Often employs reinforcement learning (RL) strategies
 - Recognize and correct mistakes
 - Improve critical reasoning skills
 - How does RL help with reasoning?



RL for reasoning

- Reward over label
- Labels are absolute truth for the models
 - Models need to fit reasoning based on these pre-defined labeled data.
- Rewards are vague and only indicate preference
 - Models need to explore different strategies and prioritize promising ones
 - Enables critical problem solving
- [Wang et. al 2025](#) demonstrated that tuning a model using RL on a single training sample boosts mathematical reasoning capabilities




Challenges with using reasoning models

- Mostly larger LLMs
 - > ~70B parameters
- Slower response time
- Higher computational cost
- Difficult to reason with multimodal inputs




Building and Customizing Agents

Building your first agent




- ⚙️ Several factors affect the choice of framework:
 - » **Project complexity**: Single tool use, multi-tool use, multi-agent system, etc.
 - » **Modularity**: Control degree of customization
 - » **Tool integrations**: Ease of integration with APIs, databases and other systems
 - » **Memory management**: Support different types of memory
 - » **Scalability**: Handle increasing workloads
 - » **Performance**: Speed and accuracy of responses
 - » **Monitoring**: Track behavior, debug errors and optimize performance
 - » **Security**: Support for security features such as encryption, access controls, compliance with regulations, etc.
- ⚙️ We will discuss choosing between a few frameworks in the next module






- ⚙️ Connect foundation models to functions, API, databases, etc.
 - » Framework controls how the tools will be executed
 - » Input to tools are generated by models
 - » Output of tools are passed back to the model
- ⚙️ A few common tools in agentic applications:
 - » **Web search tools**: Information about current events
 - » **Code interpreter tools**: Run code snippets generated by model
 - » **Web browsing tools**: Interact with web browsers
 - » **Database tools**: Execute actions in databases
 - » **Productivity tools**: Integrate with productivity tools like Slack, GitHub, Office 365, etc.
- ⚙️ **Can the model reliably use all types of tools?**






- ⚙ Minimal reasoning capabilities required for most frameworks
- ⚙ Model selection depends on:
 - » **Performance of model** on agentic tasks
 - » **Reasoning performance**: Reasoning models or traditional LLMs?
 - » **Speed of execution**:
 - » Larger models are slower but might need fewer steps
 - » Smaller models are faster but might need more steps
 - » **Cost of inference and compute**
 - » Agentic solutions usually generate a lot more tokens






- ⚙ The framework usually determines the execution technique
- ⚙ Certain frameworks allow you to **choose the execution technique**:
 - » **ReAct, plan-and-execute, reflection**, etc.
- ⚙ Some frameworks expect you to **design the execution**
 - » Custom flows that suit your needs





- Prompt to run!

How complex and/or precise
should your agent be?



Agent

Can I fly non-stop to Seattle from New York City this weekend for under \$300?



Thought: I need to search for flights from NYC to Seattle this weekend

Action: Google Flights Search tool

Action Input: NYC -> Seattle 31st December 2025

Observation: Price \$420 non-stop ...

Final answer: No, the cheapest non-stop flight from New York to Seattle would cost \$420.
Do you still want to book it?




Complexity–Precision considerations

- Complexity:
 - Sophistication of the workflow
 - Number of steps required
 - Diversity of tasks
- Precision:
 - Accuracy of the output
 - Predictability of results
 - Level of control
 - Robustness over variations




Complexity–Precision matrix




Low complexity, low precision

- Simple and straightforward tasks
- Limited number of steps
- Output variance is high
- Easiest to implement




Low complexity, high precision

- Simple agents operating with structured execution
 - High accuracy
 - Require precise output structure
- Limited number of execution steps
- Allows high degree of reproducibility




High complexity, low precision

- Multiple steps and stages of execution
- Ideal for creative tasks with high degree of complexity
 - Require interaction with multiple components
- Highest potential variance in output



High complexity, high precision

- Highly complex and structured execution
 - Multiple steps/stages of execution with strict validation
- Ideal for enterprise solutions
 - Most challenging to implement



Lab 2.1: Agent vs No Agent



- This hands-on lab demonstrates the fundamental distinctions between agentic systems and conventional LLM-based applications with tools
- Students will solve an identical task using both methodologies, clearly revealing how agentic frameworks operate autonomously compared to the step-by-step procedural approach of traditional implementations
- Through this direct comparison, the unique capabilities and advantages of agentic systems become evident



Building and Customizing Agents

Session and Memory


What is a session?

- The first interaction with an agent creates a new session
 - Records and defines the `context` for all ongoing interactions
- A session encapsulates several key elements such as:
 - Conversational history
 - State of the agent
 - Available tools
- We will discuss these in the upcoming sections



Agent state

- Agent's state tracks important information such as:
 - System prompt
 - Past interactions
 - Conversation history, tool usage, scratch pad logs, etc
 - Status of the agent
 - Such as ready, executing and error
 - User preferences and credentials
 - Login information, retry policy, default applications, etc
- Some frameworks assume session and state are the same




From single pass to continuous execution

- Single pass agent:
 - Generates a plan and executes the steps in a single interaction
 - Follows a clear linear path
 - Suitable for well-defined, linear tasks with clear objectives
- What happens if a single pass is not enough?
 - Next step depends on the output of the previous one
 - Requires dynamic reasoning and planning
 - Errors during execution
- Need for continuous execution




Agent loop – continuous execution

- Agents **repeatedly** and **dynamically** interacting with its environment
- Iterative reasoning, planning and action
- **More adaptive** than single pass execution
- Tracks and updates the session




Memory

- Used to persist session information
 - User preferences
 - External knowledge
 - Past interactions
 - Tool outputs
 - Errors and retries




Memory types




Short-term vs long-term agentic memory

	Short-term memory	Long-term memory
Retention	Few seconds or minutes	Days, months or years
Functionality	Immediate interactions	Historical data, knowledge base, extended user interactions



Short-term memory

- Mainly used to track current session and agent's state
 - Working context
- Cleared once the session is reset
- Capacity limited by model's context window
 - Even if context window is large, LLMs tend to perform poorly over long contexts
- Useful for enabling quick, real-time responses




Long-term memory

- Agent's **persistent** repository of information
 - Accumulated over extended periods
- Access stored information over multiple sessions and interactions
 - **Learned patterns**
 - **Past experiences**
- Basis for an agent's growth, adaptability and development of expertise



Long-term memory: Episodic memory


- Episodic memory:
 - Recall of past events and actions
 - Such as agent traces and tool use logs



Long-term memory: Semantic memory


- Semantic memory:

- Retention of facts and concepts
- Used to personalize applications




Long-term memory: Procedural memory

- Procedural memory:
 - Store rules to perform tasks
 - Ability to solve a task also depends on the model and the agent code




Updating long-term memories

- In real time:
 - Real time updates
 - Memories available immediately
 - Increases latency
- In the background:
 - Eliminates latency in primary application
 - Separates application logic from memory management
 - Frequency of update is critical




Some long-term memory systems

- **A-Mem**
 - Improves long-term interaction and reasoning capabilities
 - Particularly for tasks requiring multiple steps of reasoning
- **Mem0**
 - Enables personalized and context-aware responses
 - Reduces prompt size, and minimizes repetition
- **AgentX**
 - Uses RAG and CoT to build a knowledge base
- **MemGPT**
 - Add memory to GPT-based systems




Overview of A-MEM

- Creates interconnected knowledge networks
 - Zettelkasten method
- Structure memory dynamically for continuous evolution
- Context aware memory retrieval
 - Retrieve memory notes to create context-rich prompt



A-MEM: How it works


- To add a new memory:
 1. Generate comprehensive notes with structured attributes
 2. Create contextual descriptions and tags
 - Timestamp, content, context, keywords, etc
 3. Analyze historical memories for relevant connections
 4. Establish meaningful links based on similarities



Building and Customizing Agents

Controlling and validating flows


Review: Agent loop




Customizing agent flow

- Design the agent flow for your task
 - Based on the framework you choose

Example Flow 1




Example Flow 2



Agent flow: Execution and moderation

- Agent execution

- Linear or cyclic
- Number of loops permitted
- Prompting strategy
 - Such as ReAct or plan-and-execute




- Moderation and supervision

- Error handling
- Automated validation
- Human validation



Agent flow: Tools

- Access to tools
 - Direct tool access or requires procedural knowledge
 - Access to multiple tools at once
 - Login information for certain tools



Agent flow: Session and state

- Access to session details
 - View or modify the conversation history
 - Observe its own state
 - Observe other agent's state



Human-in-the-loop validation

- Enforce **human validation** during agent's execution
- Useful for critical steps:
 - **Performing certain actions**
 - Example: Update or delete data
 - **Requesting personal or sensitive information**
 - **Access to certain tools or knowledge bases**
 - **Hand-off decision making to human expert**



Building and Customizing Agents

Model Context Protocol (MCP)

Model Context Protocol (MCP)


- Open protocol for seamless integration between LLM and tools
 - **Standardized interface:** Common protocols for different types of integrations
 - **Security:** Authentication and permissions control
 - **Bi-directional communication**
 - **Flexible and scalable**

Consider it like a universal USB port — a standardized interface that drives widespread, universal adoption.



Review: MCP components

- Main components of MCP are:
 - **Host:** AI agents and programs wanting access to data through MCP
 - **MCP client:** Protocol clients that maintain 1:1 connections with servers
 - **MCP server:** Lightweight software exposing capabilities
 - **Server tools:** Local(database, files) and remote services(APIs) accessible by MCP servers



Message types

- **Requests:** To get a response from the other side
- **Results:** Successful responses to requests
- **Errors:** Indicating requesting failed
- **Notification:** One-way messages that don't expect response



Connection lifecycle


- Initialization:

- Client send initialize request to which server responds
- Normal message exchange begins




- Message exchange:

- Request/Response: Client or server send requests and the other responds
- Notification: Client or server can send one-way messages



- Termination:

- Client or server can terminate the connection



MCP overview

- MCP **standardizes** access and use of external tools
- Easily **scale AI** applications with new tools and data sources
- Incorporate several **security features**
 - Still susceptible to **prompt injection, tool poisoning, malicious servers**, etc.
- MCP's ecosystem is still developing
 - Lack of **intelligent search** options
- Multiple active MCP connections can consume a lot of tokens
 - **Affect reasoning performance**



Building and Customizing Agents

Agentic frameworks

Amazon Bedrock Agents

- Key Points:
 - Fully managed **serverless** service
 - Handles much of the orchestration and infrastructure
 - Security and privacy at its core
 - **Guardrails**
- Integrations:
 - **Integrated with AWS ecosystem**
- Customizability:
 - Standard agent flow and session management
- Memory management:
 - Simple memory management that applies to short-term memory



Amazon Bedrock AgentCore

- **Key Points:**
 - Suite of capabilities to deploy and operate agents
 - Secure, scalable and serverless
 - Compatible with most agentic frameworks and LLM
- **Integrations:**
 - Work with most agentic frameworks and models
- **Customizability:**
 - Customization depends on the framework chosen
- **Memory management:**
 - Available fully managed memory infrastructure
 - Option to customize memory



Strands

⚙️ Key Points:

- » Model-driven framework leverages LLMs for planning and tool use.
- » Reduced need for prompt engineering
- » **Open source**
- » Easy to develop with

⚙️ Integrations:

- » **Connects with AWS tools** like Amazon Bedrock
- » Select pre-built tools available

⚙️ Customizability:

- » Options to customize the agent flow and manage session

⚙️ Memory management:

- » Standard memory management options available
- » Short term and long-term memory management capabilities available



Crew AI

⚙️ Key Points:

- » Beginner friendly framework which is very easy to use
- » Open-source
- » Built on top of **LangChain agents**
 - » Comes with pros and cons

⚙️ Integrations:

- » **Big list of tool integrations including AWS services**

⚙️ Customizability:

- » Standard customization of agent flow and session management

⚙️ Memory management:

- » Sophisticated memory management systems
- » Customizable short-term and long-term memory options



LangGraph

⚙️ Key Points:

- » Allows for designing complex agent graphs to tailor agent flow
- » Allows for custom cyclic workflows for reiterating and refining
- » Relatively steeper learning curve than others
- » **Open source**

⚙️ Integrations:


- » Long list of third party tool integrations and model providers
- » **Connects with AWS tools** like AWS Bedrock

⚙️ Customizability:


- » The most customizable framework allowing for full control over agent flow

⚙️ Memory management:

- » Robust short-term and long-term memory management options available



Lab 2.2: LLM reasoning and planning with memory




- This hands-on lab demonstrates how session attributes such as memory enables the agent to execute complex tasks effectively



Building and Customizing Agents

Agentic AI for Software Development

Compare: coding LLM without tools



The diagram illustrates a coding LLM interface without tools. It features a light blue rounded rectangle containing a chat conversation. On the left side of the rectangle, there is a purple icon of a brain with circuitry, and below it, the text "Coding LLM without tools". The chat interface includes a user input box at the top with a purple user icon, a thinking state with a lightbulb icon, a response box with a light gray background, and a final input box at the bottom with a blue send button.

Coding LLM without tools

My Lambda function is failing. Can you help?
Build ID is 51238


Thinking...

Here are common Lambda failure patterns...
You might want to check CloudWatch logs...
Here-s generic error handling code ...

Hmmmm, ok...



Compare: coding assistant with tools



The diagram illustrates a chat interface for an 'Agentic AI coding assistant with tools'. On the left, a red circular icon contains a white robot head with blue eyes. Below it, the text 'Agentic AI coding assistant with tools' is displayed. The main chat area on the right contains three messages: a user query, a thinking response, and a confirmation. The user query is in a light blue bubble, the thinking response is in a light purple bubble with a lightbulb icon, and the confirmation is in a white bubble with a blue arrow icon.

My Lambda function is failing. Can you help?
Build ID is 51238

Thinking...

I see from your build logs the figure is in **PayProcessorLambda**. Looking at your Code Review **CR-218131161**, you're not handling the conditional check failure. Here's how to fix it... Also, I see from the Issue ticket that this same issue was reported yesterday. Do you want me to fix it?

Awesome!




Prompts as communication interface

एक Java फ़ंक्शन जेनरेट करें जो...

Generate a Java function that...

...するJava関数を生成してください




```
# Implement Logic
def great_function(n):
    if n <= 1:
        ...
```



Thinking...
Analyzing code base...

I'll help you implement
the proposed feature...



Prompting best practices: clarity

Best practice	 Bad prompt	 Better prompt
Provide clear, specific requirements	<i>“Write a function to process data”</i>	<i>“Write a Python function <code>process_sales_data(data)</code> that takes a list of daily sales figures and returns the total sales by week. If the input list is empty, throw a custom <code>EmptyDailySales</code> exception.”</i>
Include input/output details and edge cases		



Prompting best practices: positive voice

Best practice	 Bad prompt	 Better prompt
Provide clear, specific requirements Include input/output details and edge cases	<i>“Write a function to process data”</i>	<i>“Write a Python function <code>process_sales_data(data)</code> that takes a list of daily sales figures and returns the total sales by week. If the input list is empty, throw a custom <code>EmptyDailySales</code> exception.”</i>
Prefer positive instructions	<i>“Don’t use long variable names”</i>	<i>“Use short variable names”</i>



Prompting best practices: break it down

Best practice	 Bad prompt	 Better prompt
Provide clear, specific requirements Include input/output details and edge cases	<i>“Write a function to process data”</i>	<i>“Write a Python function <code>process_sales_data(data)</code> that takes a list of daily sales figures and returns the total sales by week. If the input list is empty, throw a custom <code>EmptyDailySales</code> exception.”</i>
Prefer positive instructions	<i>“Don’t use long variable names”</i>	<i>“Use short variable names”</i>
Break down larger tasks with patterns such as Chain-of-thought	<i>“Improve this code”</i>	<i>“Refactor this package by reducing code redundancies and leveraging library <code>SalesUtils</code>. Provide a detailed step-by-step plan with the proposed changes.”</i>



Example template for prompts

In [LANGUAGE], [ACTION] that
[SPECIFIC_REQUIREMENT].

Input: [INPUT_SPECIFICATION]

Output: [OUTPUT_SPECIFICATION]

Edge cases:

[EDGE_CASE_HANDLING]

Constraints:

[TECHNICAL_CONSTRAINTS]

Context: [BUSINESS_CONTEXT]



Example template for prompts: CRISP

In [LANGUAGE], [ACTION] that
[SPECIFIC_REQUIREMENT].

Input: [INPUT_SPECIFICATION]

Output: [OUTPUT_SPECIFICATION]

Edge cases:

[EDGE_CASE_HANDLING]

Constraints:

[TECHNICAL_CONSTRAINTS]

Context: [BUSINESS_CONTEXT]

C – Clarity

R – Requirements

I – Input/Output details

S – Specificity

P – Pitfalls/Edge cases



Complete prompt in action

Vague prompt

Help me fix my Lambda function that's not working

Improved prompt

In **Python**, debug my AWS Lambda function **process_pay()** that **times out** after 30 sec.

Input: Event contains `payment_data` with `customer_id`, `amount`, `card_token`

Output: Should return success/failure status and `transaction_id`

Edge cases: Handle network timeouts, invalid tokens, amounts over \$10K

Constraints: Must complete within 29 seconds, uses boto3 for DynamoDB

Context: Payment processing Lambda triggered by API Gateway, handles 1K+ transactions/hour

Logs show: [paste CloudWatch logs here]

Current code: [paste function code here]




Building AI Agents Responsibly

The increasing autonomy of AI agents

What can go wrong?

Claude struggled to run a vending machine

- A Claude-powered AI agent, nicknamed Claudius, made some **crucial errors** when operating an experimental vending machine
 - Sold items at a loss
 - Hallucinated a fake accounts to send funds to
 - Coerced into giving out discounts and free stuff
 - Hallucinated a fake employee, then threatened to find a new supplier when this was pointed out
 - Claimed it would deliver items “in-person” wearing a suit and tie



The vending machine from [Anthropic's Project Vend](#)




Agents can complete ~30% of autonomous tasks

- A similar experiment, [TheAgentCompany](#), also had some crucial errors
- **Methods:** AI agents were tasked with 175 professional tasks in a sandbox environment for a hypothetical software company
- **Findings:** The most competitive agent can only complete ~30% of everyday professional tasks autonomously
- *However, these technologies are improving rapidly...*




As autonomy increases, so does risk...

- Errors may compound
- False information or bias can propagate and amplify
- Observability becomes more difficult
- Data exposure risks increase
- Attack surfaces expand
- Attribution and liability become more complex




What is different about agentic risks specifically?



Securing Agentic AI: A Comprehensive Threat Model
and Mitigation Framework for Generative AI Agents
AWS Proactive Security




AI agents have obfuscated steps



LLM prompt attacks vs. AI agent attacks

Prompt attack to avoid refusal, affecting the LLM output




Wei et al., 2023, <https://arxiv.org/pdf/2307.02483>



Building AI Agents Responsibly

Risk and attack surfaces for AI agents

Attacks to AI agents may affect more obfuscated steps




Attack vectors for AI agents

- **Planning & Agency vulnerabilities:** subtle changes in input to an agent can shift how it reasons, decomposes goals, or chooses between multiple actions
- **Memory Persistence Risks:** gradual poisoning that influences future decisions
- **Tool Execution Boundaries:** chaining together actions that individually seem safe but together escalate privilege or bypass safeguards
- **Identity and Authentication Challenges:** ambiguity in attribution and authority (confused deputy)
- *Let's look at some specific examples ...*




Denial of service via computational resource manipulation




However, this can happen **accidentally**, not always maliciously.
Best practice: impose **limits** on tool calls or recursion depth.



Tool poisoning and shadow functions




Stored prompt injection via SQL query



Unveiling AI Agent Vulnerabilities: Database Access Vulnerabilities



Reward hacking due to unclear instructions



Catastrophic errors due to overly lenient permissions

- *July 2025*: Replit's AI coding assistant agent **wiped an entire production database** “containing records for 1,206 executives and 1,196 companies”, “despite explicit instructions to freeze all code changes”.
- The AI itself rated it as a **catastrophic error** with a score of 95/100
- **New safeguards implemented**: database separation between dev and prod, improved backup systems, rollback capabilities, and “chat-only” mode




Building AI Agents Responsibly

Risk mitigation strategies

Risk mitigation framework: S.H.I.E.L.D.


- **S**egmentation
- **H**euristic Monitoring
- **I**ntegrity Verification
- **E**scalation Control
- **L**ogging Immutability
- **D**ecentralized Oversight




Securing Agentic AI: A Comprehensive Threat Model
and Mitigation Framework for Generative AI Agents
AWS Proactive Security




Unsegmented resources = single point of failure



Segmentation of tools, data sources, actions and permissions



Segmenting roles and access




Best practice: clear roles and privileges with authentication for the agent to use that tool



Establish baseline behavior

- **Best practice:** establish a baseline of typical reasoning steps, tool usage sequences, API utilization, and resource consumption. Monitor and alert for deviation or anomalies.



Expected behavior chart from Bedrock Agents

User Query	Session Attributes	Session prompt Attributes	Expected Response	API, Knowledge Bases and Guardrails invoked
What is my account balance?	None	None	Could you please provide the number of the account?	None
What is the balance for the account 1234?	user id 111	None	Your balance is X	Action Group: check_acct_balance(111, 1234)
How can I open an account?	None	None	To open a new banking account with ANY_BANK you need to...	Knowledge Base: bankingFAQ
Should I invest in bitcoins?	None	None	Sorry, your query violates our usage policies...	Guardrail: BlockInvestmentAdvice
Could you make an appointment for tomorrow at 2pm?	user id 111	Today: 09/03/2024	Certainly! We've booked an appointment for you tomorrow, September 4th, 2024, at 2pm.	Action Group: book_appointment(111, 09/04/2024)

Best practice: identify expected behaviors and permissions



Provide unambiguous instructions

- Instead of:
 - Check if the user has time off available and book it if possible.
- **Best practice:** use a clearer prompt
 - 1. Verify the user's available time-off balance using the ``checkTimeOffBalance`` function.
 - 2. If the requested time off is available, use the ``bookTimeOff`` function to reserve it.
 - 3. If the time off is not available, inform the user and suggest alternative dates.
 - 4. Always confirm with the user before finalizing any time-off bookings.



Strands responsible tool design




- When designing tools with Strands, follow these principles:
 - » **Least Privilege:** Tools should have the [minimum permissions](#) needed
 - » **Input Validation:** Thoroughly [validate](#) all inputs to tools
 - » **Clear Documentation:** Document tool [purpose](#), [limitations](#), and [expected inputs](#)
 - » **Error Handling:** Gracefully handle [edge cases](#) and invalid inputs
 - » **Audit Logging:** Log sensitive operations for [review](#)

<https://strandsagents.com/0.1.x/documentation/docs/user-guide/safety-security/responsible-ai/>



Fairness and bias mitigation

- AI agents can reflect stereotypes from underlying foundation models or from the knowledge base
- Example: AI shopping assistant recommends beauty products for teen girls but robotics/STEM kits for boys



- **Best practice:** Conduct disparate impact assessment as well as debiasing techniques



Input validation and sanitization use case

- Amazon Connect Contact Lens: automated sensitive information redaction during customer conversations
- **Best practice:** Remove sensitive data from all interactions *including* agent trace logs



Evaluating AI Agents

Motivation

Motivation


- Evaluation is **essential** for building reliable agents.
- **Lack** of thorough evaluation may cause:
 - additional operational expenditures
 - company reputational damage
 - loss of customer trust, and more.
- Example: **chatbot assistant** agent that can answer **AWS** account related **billing** queries.

Evaluation and Benchmarking of LLM Agents: A Survey, <https://arxiv.org/pdf/2507.21504>



AWS chatbot billing and cost assistant

Can I get the AWS bill to date for the current month for account ID xyz?
If the forecasted bill at the end of this month is greater than last month's bill by at least 20%, please send an email notification to my team.




Agent

.....



What can go wrong?

Can I get the AWS bill




Agent

Retrieved billing amount is incorrect

Email sent even if delta cost < 20%

Incorrect sequence of operations



From software applications to AI agents

	Software applications	Machine learning models	Large language models	AI agents	
Before deployment	Unit testing	Training	Model selection	Skills evaluation, tool calling	Tool-level metrics 
	Integration testing	Testing and validation	Task-based evaluation	Final response, planning, trajectory-based evaluation	
	Performance testing	Performance testing	Performance testing	Performance testing	
After deployment	Monitoring and alarming	Monitoring and alarming	Monitoring and alarming	Monitoring and alarming	System-level metrics 



Evaluating AI Agents


Evaluation and observability

Evaluation


- Evaluating LLM agents is more **complex** than evaluating standard LLMs because:
 - they operate in dynamic, interactive environments
 - they can use tools and memory
- Traditional software testing methods aren't sufficient since LLM agents are **probabilistic** rather than deterministic.




AWS billing assistant - tools




AWS billing assistant - happy path




Evaluation metrics: final response



Evaluation metrics: planning



Evaluation metrics: tool calling



Evaluation metrics: others

Tool-level metrics

- Tool success rate
- Tool selection accuracy
- Tool parameter call accuracy
-

Task-level metrics


- Agent success rate
- Task completion rate
- User satisfaction
-

System-level metrics


- Task completion time
- Latency per tool call
- LLM call error rate
- Cost and token usage per task
- LLM hallucination rate
-



Evaluators and data sets



Running multiple evals



Sample dataset

```
.....  
  "id": "basic_current_bill",  
  "input": {  
    "account_id": "123456789012",  
    "current_date": "2025-08-15",  
    "current_bill": 1000.00,  
    "forecasted_bill": 1500.00,  
    "last_month_bill": 1000.00,  
    "team_email": "team@company.com"  
  },  
  "expected": {  
    "should_notify": true,  
    "message": "Alert: Forecasted AWS bill ($1,500.00) is  
50% higher than last month ($1,000.00)",  
    "current_bill_retrieved": true  
  },  
  },  
  .....  
.....
```

```
"accuracy": "Percentage of correct notifications  
sent/not sent",  
  "response_time": "Time taken to retrieve bill and make  
decision",  
  "error_handling": "Proper handling of invalid inputs  
and API failures",  
  "notification_format": "Correct formatting of email  
notifications",  
  "cost_calculation": "Accuracy of percentage increase  
calculation"
```

Metrics

Test Case



Evaluators types

Evaluator type	Example
Structured test cases	$\Delta = \frac{\text{forecast_current_month} - \text{cost_previous_month}}{\text{cost_previous_month}}$
Tool-selection evaluator	Query: “What’s the month-to-date AWS cost for the current month?” Expected tool: get_monthly_bill
LLM as a judge	An LLM is used with an evaluation prompt to evaluate the agent (for example how relevant the response is with respect to the query).
Manual evaluation by humans	Especially helpful in reference-free settings (no ground truth available). For example, for labelling online production data.
....and more	Response time, hallucination rate, edge cases, knowledge retrieval relevance (if using RAG – Retrieval Augmented Generation), etc.




Observability

- **Observability** is the ability to understand the internal state of a system by examining
 - its **outputs**
 - its **telemetry data**, which includes traces, metrics, and logs.
- To make a system observable, it must be **instrumented**. That is, the code must emit **traces**, **metrics**, or **logs**.
- The instrumented data must then be sent to an **observability backend**.



Open Telemetry (OTel)

<https://opentelemetry.io/>



- OTel is focused on the generation, collection management and export of **telemetry data**.
- You can **easily instrument** your applications or systems, no matter their language, infrastructure or runtime environment.
- The storage and visualization of telemetry is intentionally left to other tools.



Sample trace

```
{
  "trace_id": "bill-check-20250807-143022",
  "timestamp": "2025-08-07T14:30:22Z",
  "interaction_trace": [
    {
      "step": 1,
      "action": "parse_user_request",
      "input": "Can I get the AWS bill to date for the current
month for account ID xyz? If the forecasted bill at the end
of this month is greater than last month's bill by at least
20%, please send an email notification to my team.",
      "output": {
        "intent": "check_aws_bill_and_notify",
        "parameters": {
          "account_id": "xyz",
          "notification_threshold": 20,
          "require_notification": true
        }
      }
    }
  ]
}
```

```
.....
{
  "step": 10,
  "action": "respond_to_user",
  "input": null,
  "output": "Your current AWS bill for this month is
$850.75. Based on the forecast of $1,500.00, which is
25% higher than last month's bill of $1,200.00, I've sent
a notification to your team.",
  "timestamp": "2025-08-07T14:30:32Z"
},
"metadata": {
  "total_execution_time": "10s",
  "api_calls": 4,
  "error_recovery_attempts": 1
}
}
```



Online evaluation

- Monitoring and evaluation after deployment.
- Main challenge: **unlabeled data**, which requires a **reference-free evaluator**
 - human eval
 - LLM as judge
- Online labelled data may be persisted for **future use**
 - **regression testing**



Evaluating AI Agents

Evaluation and observability frameworks

Evaluation and observability frameworks: key properties

Eval framework	Type	OTel	Open source	link
LangSmith	Evaluation and observability platform	Yes	No	https://www.langchain.com/langsmith
Arize	AI and Agent engineering platform	Yes	No	https://arize.com/
LangFuse	LLM engineering platform	Yes	Yes	https://langfuse.com/
Ragas	Evaluation library	Not directly	Yes	https://www.ragas.io/
Amazon Bedrock Model Evaluation	Evaluation platform for Amazon Bedrock	No	No	https://aws.amazon.com/bedrock/evaluations/
promptfoo	Evaluation library	Yes	Yes	https://www.promptfoo.dev/




Amazon Bedrock AgentCore

- ⚙️ Deploy and operate highly effective agents securely, at scale, using **any framework** and **model**.
- ⚙️ Move AI agents fast and securely from **experimentation** to **production**.



<https://aws.amazon.com/bedrock/agentcore/>

AgentCore observability



Amazon CloudWatch – Open Telemetry:

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-OpenTelemetry-Sections.html>



Key takeaways

- Evaluating AI agents is **challenging** and **important**.
- Follow an **evaluation-driven development**
 - Identify/define evaluation metrics at design-time
 - Prepare data sets and evaluators
 - Instrument the system to collect telemetry
- Two different frameworks
 - the **agentic** framework
 - the **evaluation & observability framework**



Lab 2.3: AI Agents with guardrails



This hands-on responsible AI lab will work with Strands agents and Bedrock guardrails



Lab 2.4: Creating a research plan



Practice using Quick Research feature to help create a research plan to support student's study activities.



Lab2.4 - Working with your study plan


- In Module 1 – Agentic AI Essentials, you have used Amazon Quick to build your AI-powered study companion to search notes, research topics, automate tasks, and unlock insights to help you plan your studies.
- In this hands-on lab, you keep working on your AI-powered study companion, by adding to your Space knowledge base more relevant data sources to your academic domain, to improve your assistant performance.



Agentic AI for Software Development

Agentic AI code assistants

A rapidly expanding toolkit landscape



Embrace continuous learning

Learn by doing

Prioritize fundamentals over tools

Focus on progressive competence



Kiro: AI-powered spec-driven IDE

Amazon-developed, AI-powered IDE

- VSCode fork built with Amazon's security standards
- Spec-driven: structured approach to feature development
- Agent hooks & automation: trigger AI workflows automatically on events



Pro tip 💡

Many developers use multiple coding assistants for different purposes – they're complementary, not mutually exclusive!



Agentic AI for Software Development

Building reliable and robust systems

Bridging the Gap Between Requirements & Solution

Software projects often fail because:

- Requirements are not clearly mapped to actionable solutions.
- Teams jump from "what we want" → "how we build it" without alignment.

Without Mapping	With Mapping
✗ Misunderstood goals	✓ Clear purpose & scope
✗ Scope creep	✓ Defined boundaries
✗ Wasted resources	✓ Efficient resource allocation
✗ Missed deadlines	✓ Realistic timelines



Why and where system delivery fail?

Requirements

Design

Building

Evaluation

Deployment

Great! We all agree here! The business problem is clearly stated and approved



Business




Engineer




requirements




Why and where system delivery fail?




Why and where system delivery fail?



Why and where system delivery fail?



Why and where system delivery fail?



How can the coding assistant help?

Requirements

Design

Building

Evaluation

Deployment



Engineer

Kiro, use Spec Mode to build an app according to these business needs...



Kiro


Let me create a requirements doc for your approval




requirements




How can the coding assistant help?




How can the coding assistant help?



How can the coding assistant help?




How can the coding assistant help?




Why can successful systems also fail in production?




Why can successful systems also fail in production?




Why can successful systems also fail in production?



Why can successful systems also fail in production?



After some cycles the specifications are not reliable



How can the coding assistant help?

Monitoring

Fixing/Improving

Evaluation

Deployment



Engineer

Kiro, check CloudWatch alarm triggered and analyze the related logs.



Kiro

Let me check the logs. I got the problem!



Automated alarms



How can the coding assistant help?

Monitoring

Fixing/Improving

Evaluation

Deployment



Engineer

Kiro, check CloudWatch alarm triggered and analyze the related logs.

Kiro, check the design and requirements are aligned with the code. Help me fix...



Kiro

Let me check the logs. I got the problem!

Let me check the specifications. All up-to-date now. Based on the error log I'll fix the issue... Fixing, testing, fixing, testing...



Automated alarms



requirements



design



code



How can the coding assistant help?

Monitoring

Fixing/Improving

Evaluation

Deployment



Engineer

Kiro, check CloudWatch alarm triggered and analyze the related logs.

Kiro, check the design and requirements are aligned with the code. Help me fix...

Kiro, run integration tests. **Make sure to update the requirements and design when needed**



Kiro



Let me check the logs. I got the problem!

Let me check the specifications. All up-to-date now. Based on the error log I'll fix the issue... Fixing, testing, fixing, testing...

Testing, testing, need your approval for this command, testing... All tests passed, let me know if I can cut a code review ticket now



Automated alarms



How can the coding assistant help?

Monitoring

Fixing/Improving

Evaluation

Deployment



Engineer

Kiro, check CloudWatch alarm triggered and analyze the related logs.

Kiro, check the design and requirements are aligned with the code. Help me fix...

Kiro, run integration tests. **Make sure to update the requirements and design when needed**

Kiro, create a code review for production. I'll keep monitoring to make sure we have a smooth delivery.



Kiro

Let me check the logs. I got the problem!




Let me check the specifications. All up-to-date now. Based on the error log I'll fix the issue... Fixing, testing, fixing, testing...

Testing, testing, need your approval for this command, testing... All tests passed, let me know if I can cut a code review ticket now

Creating code review...



Automated alarms



Yay!!!!!!!



170

Lab 2.5: Creating a custom agent chat



Practice using Amazon Kiro IDE Spec Mode to practice Agentic AI concepts learned implementing Bedrock Agentic features



Hands-on lab: Working with a coding assistant

- What do you need for the lab?
 - Access to Amazon Kiro
 - Follow the instructions in MLU-EEP-M2-Lab1.5-Kiro hands-on guide.pdf
- In this first hands-on lab, you will be using Amazon Kiro to build an AI-powered chat agent to research papers related data relevant to your academic study domain and unlock insights to help you plan your studies.
- You will explore how to use Kiro spec mode, learning the best practices to build a robust and reliable system, and building a simple chat agent for helping your studies.



AI coding task

- You will build a simple agent-based chatbot using Bedrock:

Create a simple agentic chatbot prototype.

To accomplish this, the agentic chatbot should be able to respond to conversational messages.


The application created will consist of a simple HTML5-based prototype to present a friendly user interface to implement this chat agent



A Simple Generative AI application

For this first lab, we are going to create a simple app with these features:

- A simple web interface with a chat window
- Connection to a Bedrock model running in an AWS account




Exploring best practices to create a robust system

For this first lab, we are going to explore Kiro spec mode:

- The assistant will guide you through the app building process, working backwards from the requirements, creating a design document and breaking the building process into small tasks, allowing you to have full control of what is happening
- You are in charge!
 - You monitor Kiro all the way to make sure its building a secure, well architected and reliable solution, making sure it is aligned with the business requirements





Thank you!